CAP5610: Machine Learning Final Project Presentation

Would Your Favorite Song Make it to Spotify's "Top 50 - Global" Playlist?

April 19, 2022

Bailey LaRea Nadine Rose Suma Marri

Agenda

Initial Project Questions
Description of the Data
Data Preparation
Visualizations for EDA
Revised Project Question
Baseline Model
Preliminary ML models
Next Steps

Initial Project Questions

Idea: Can we use Machine Learning techniques to help **predict** what songs will make it into Spotify's "Top 50 - Global" playlist?

Can specific track attributes be used to predict overall song popularity?

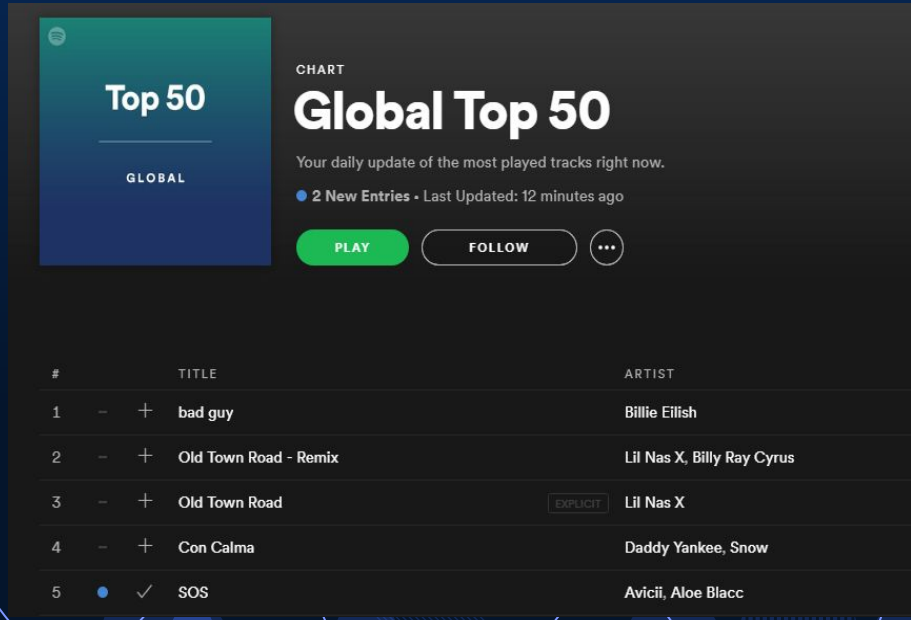

Description of the Data

There are 13 Track Attributes

TrackName
ArtistName
Genre
Beats Per Minute
Energy
Danceability
Loudness (dB)
Liveness
Valence
Length
Acousticness
Speechiness
Popularity

The Datasets

Top 50 Songs in 2019

- This dataset contains information for the top 50 songs in 2019.
- 50 rows of data

2019 Songs

- This dataset contains songs that were released in 2019.
- ≈ 11,000 rows of data

Both datasets are sourced from Kaggle

Data Preparation

Kaggle

- We used this for our initial dataset and for the top 50 songs.

Spotify for Developers - API

- The Spotify for Developers allowed us to extract and collect the data from Spotify
- We used the Client Keys to access the data on Spotify.
- After extracting the data, we had deleted and duplicate values and any records that had null values.

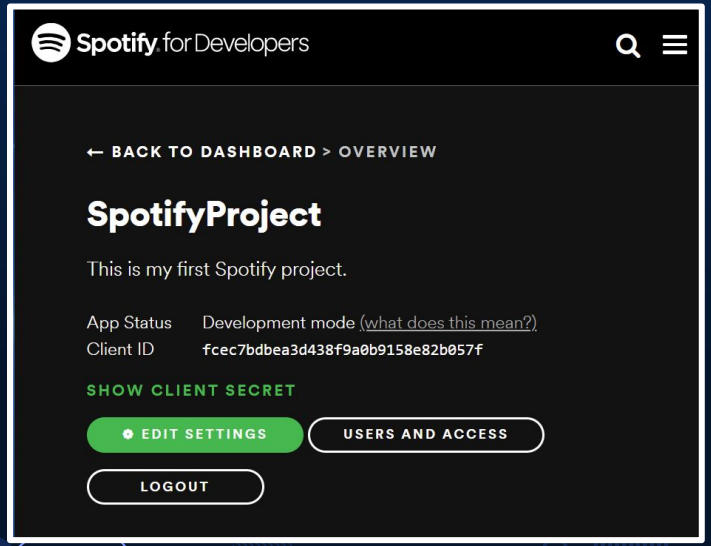

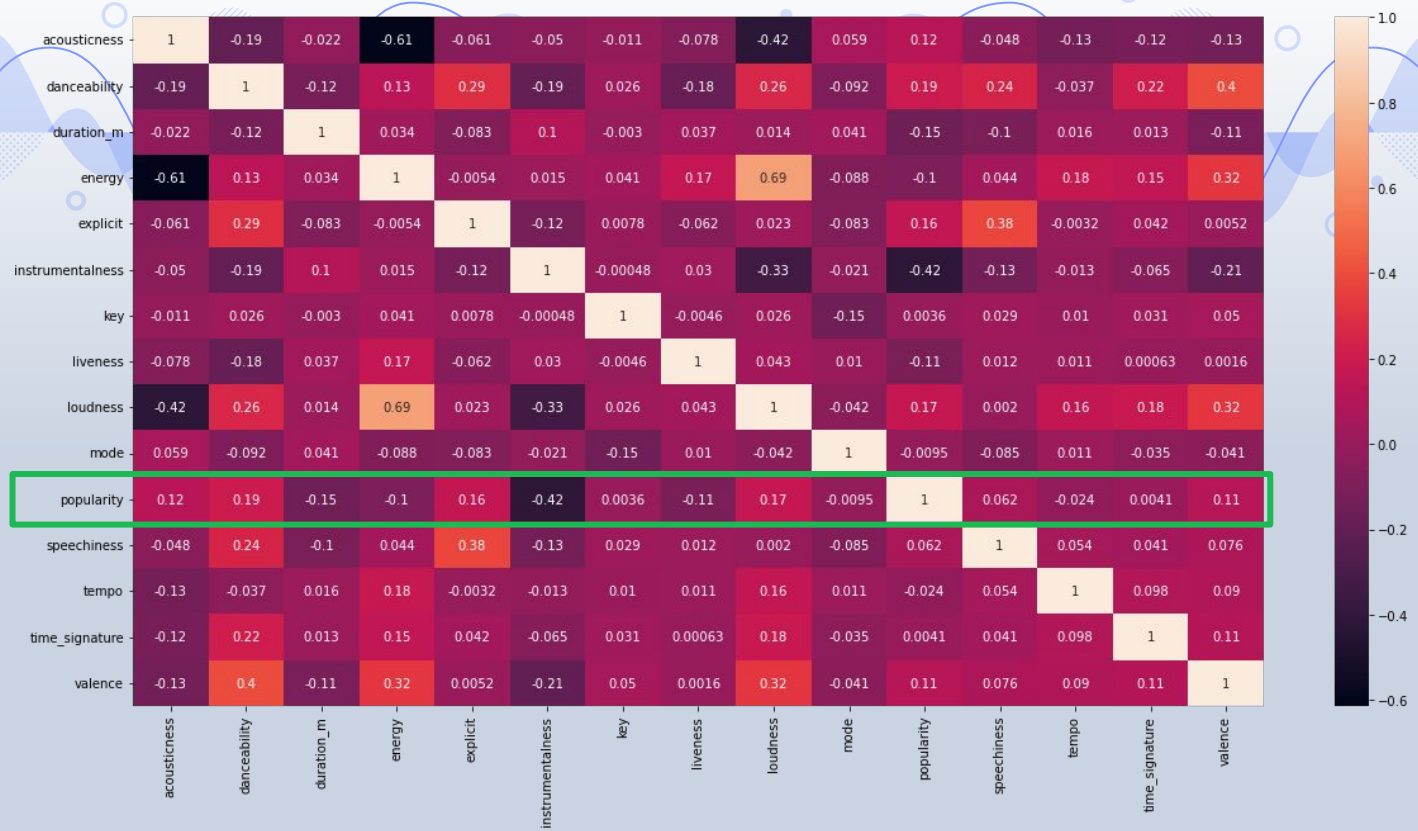

Visualizations - Heatmap

Visualizations - Scatterplots

Instrumentalness vs. Popularity

We can see that there is a **negative correlation** between both variables.

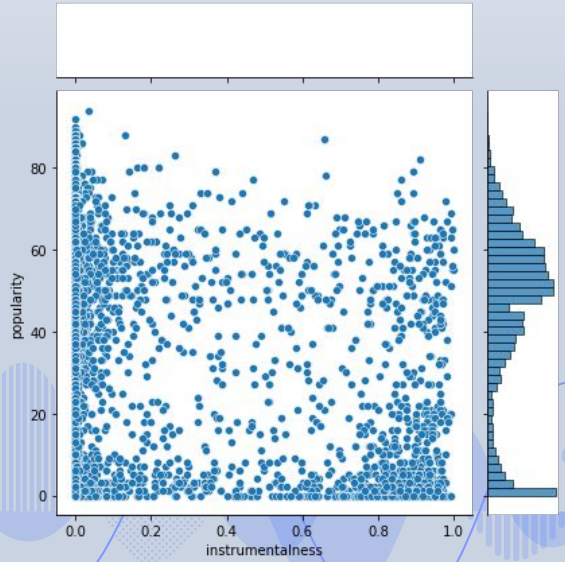

Speechiness vs. Popularity

We can see that there is **no correlation** between both variables.

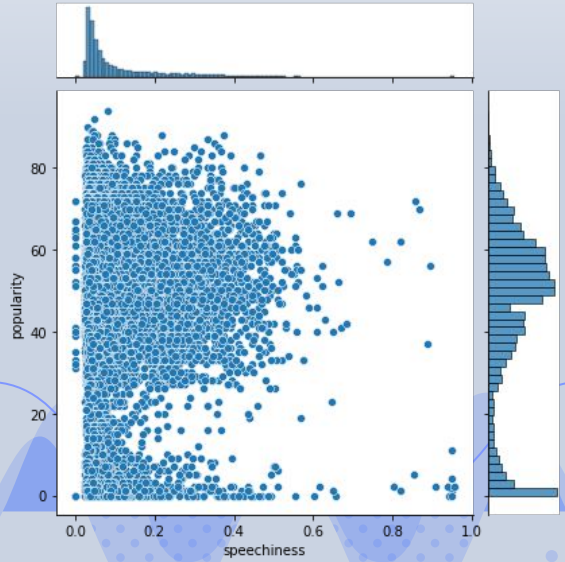

Loudness vs. Popularity

We can see that there is a **positive correlation** between both variables.

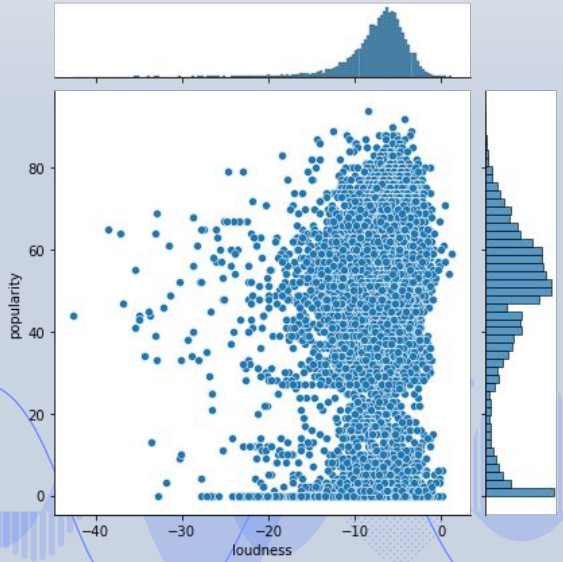

Visualizations - Histogram

Here are the following distributions of Speechiness and Loudness from another dataset. They follow distributions that align with the prior slides.

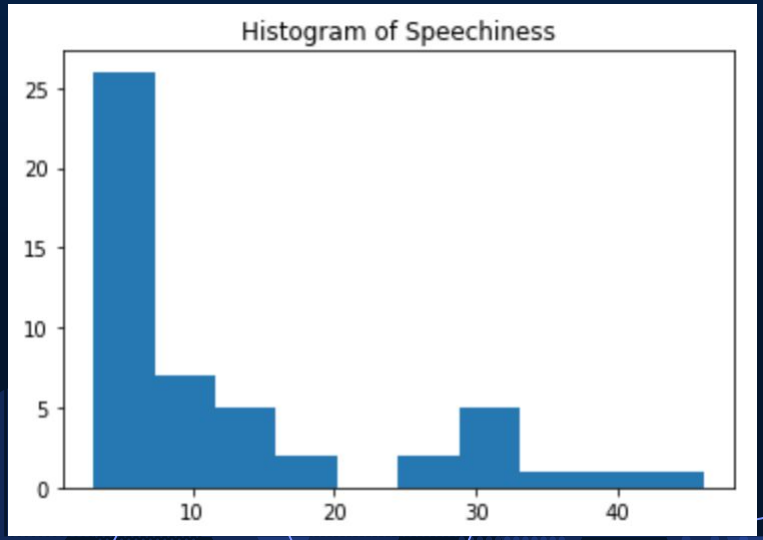

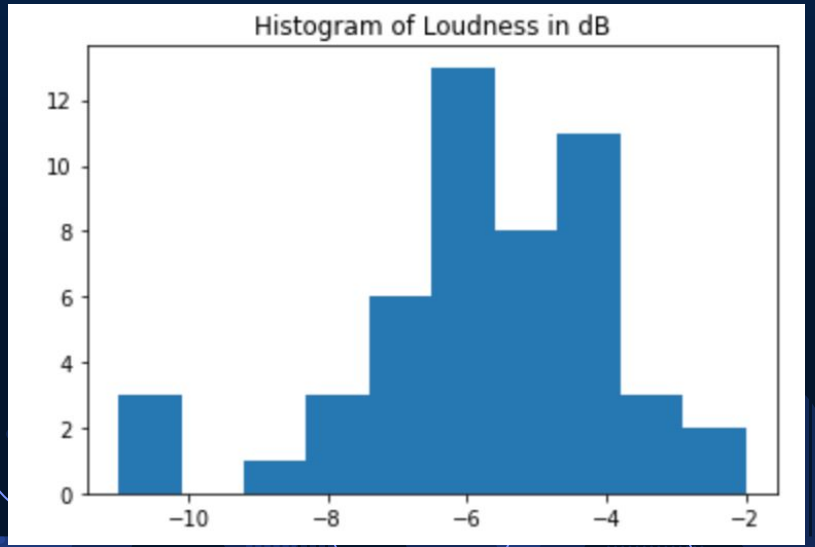

Revised Project Question

Forecasting → Time Series for Baseline Model →

- ✖ Data following sequential order with respect to time
- ✖ Time as the independent variable

Classification →

- ✅ Allows us to predict if a song will be classified as a song within the 2019 Top 50 - Global Playlist.

Baseline Model: Linear Regression

```
from sklearn.model_selection import cross_val_score
from sklearn import datasets, linear_model

scores = cross_val_score(model,X,Y, cv=5)
print("Print all scores: ", scores)
print("Mean Accuracy: ", scores.mean())
```

```
Print all scores:  [0.21003084 0.16980323 0.08999494 0.15596551 0.07952351]
Mean Accuracy:  0.14106360537917992
```

```
R2 Score:  0.2059066468330093
MAE:  13.970961045076779
MSE:  342.95193847341307
RMSE: 18.518961592740915
```

With a mean accuracy score of only ~14%, some changes need to be made!

Preliminary Version of ML Models

Decision Tree - Plot of RMSE

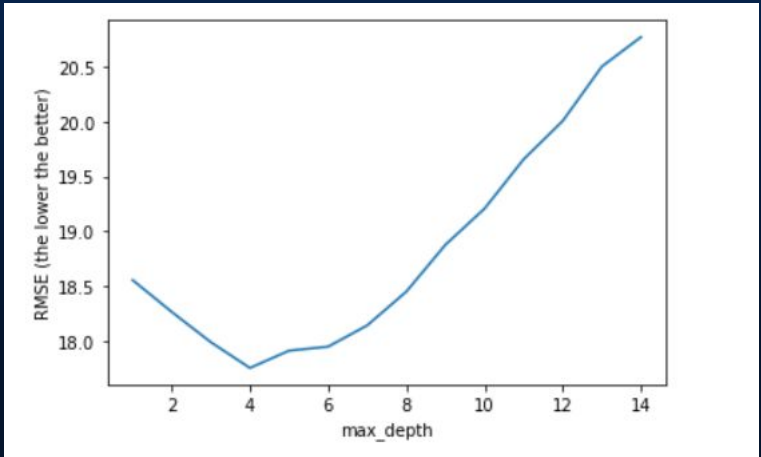

	feature	importance
4	instrumentalness	0.408603
2	duration_ms	0.128069
0	acousticness	0.120920
7	loudness	0.065383
1	danceability	0.060481
3	energy	0.051191
9	valence	0.047880
8	tempo	0.046519
6	speechiness	0.043091
5	liveness	0.027864

There is a little bit of improvement.... However, there must be a **better model** out there?

Random Forest & KNN are also used!

Random Forest: Classification

Random Forest

```
Mean Accuracy: 0.697
```

Baseline: LR

```
RMSE: 18.518961592740915
```

	feature	importance
4	instrumentalness	0.202412
0	acousticness	0.126836
2	duration_ms	0.116127
7	loudness	0.093867
3	energy	0.088406
1	danceability	0.079084
9	valence	0.076297
5	liveness	0.073639
8	tempo	0.073136
6	speechiness	0.070195

There is great improvement.... However, is there a **better model** out there?

K Nearest Neighbor (KNN)

KNN

```
Mean Accuracy Score:  0.9917773561037319
```

Baseline: LR

```
Mean Accuracy:  0.14106360537917992
```

There is significant improvement.... However, there is the risk of **overfitting**!

Next Steps

- Investigate potential overfitting of KNN model
- Use the finalized ML models to see if successful classification is achieved
- Data Preparation

Questions?

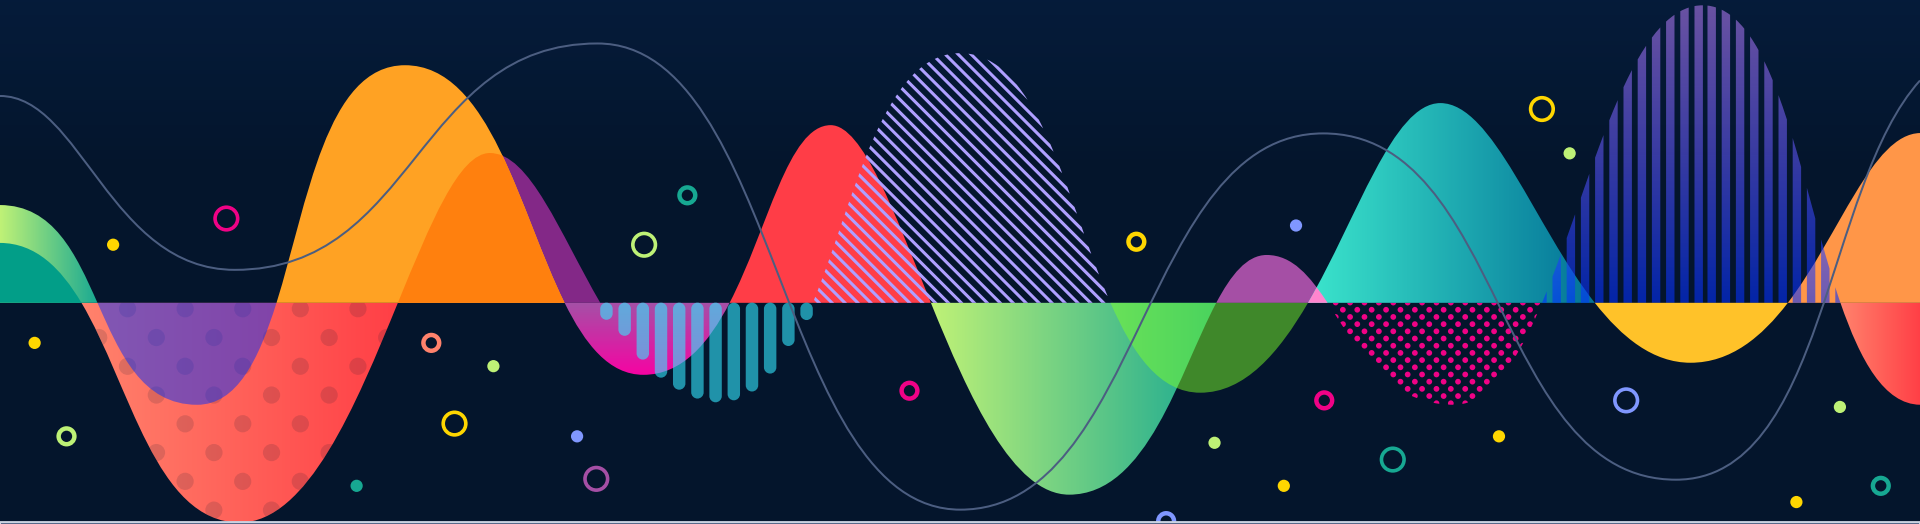

References

- https://www.kaggle.com/theoverman/the-spotify-hit-predictor-dataset?select=dataset-of-00s.csv
- https://www.kaggle.com/leonardopena/top50spotify2019
- https://towardsdatascience.com/predicting-popularity-on-spotify-when-data-needs-culture-more-than-culture-needs-data-2ed3661f75f1

Notes from Dr. Yousefi

- RMSE is used for regression models
- Mean Accuracy is used for classification models
- Need to use more regression models to compare with the baseline of linear regression
- Set a baseline for classification models and have at least 2 other models to compare to baseline
- For overfitting, it is suggested to use diagnostic plots
- All learning models should have at least 0.5 mean accuracy, check linear model again
- Decision Tree is defined as regression model in our code not classification
- Good to give the statistics of the data in the report
- 94% is good and usually not overfitting
- Regression = predict
- Argue why we are choosing the features in the analysis